

## Step 1: Importing Python Packages

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

## Step 2: Import Data

Here we are importing a real-world open sourced dataset to use for our assignment. Import your data here.

Note the data can be found at the following URL: [https://raw.githubusercontent.com/salexyun/Michener-AI-for-Clinician-Champions/main/medical\\_cost.csv](https://raw.githubusercontent.com/salexyun/Michener-AI-for-Clinician-Champions/main/medical_cost.csv)

Write your Step 2 code here in this cell

```
In [2]: URL = 'https://raw.githubusercontent.com/salexyun/Michener-AI-for-Clinician-Champions/main/medical_cost.csv'
data = pd.read_csv(URL)
data.head()
```

Out[2]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

## Step 3: Determine the dimenionality of the dataset.

Print out the total number of patients in the dataset. Print out the different unique freatures that are available for each patient.

```
In [3]: ## Write your code for Step 3 in this cell
data.describe()
```

Out[3]:

	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

```
In [4]: data['smoker'].value_counts()
```

```
Out[4]: no      1064
yes       274
Name: smoker, dtype: int64
```

```
In [5]: data['sex'].value_counts()
```

```
Out[5]: male      676
female    662
Name: sex, dtype: int64
```

```
In [6]: data['children'].value_counts()
```

```
Out[6]: 0      574
1      324
2      240
3      157
4       25
5       18
Name: children, dtype: int64
```

```
In [7]: missing_values=data.isnull().sum()
missing_values
```

```
Out[7]: age      0
sex      0
bmi      0
children 0
smoker   0
region   0
charges  0
dtype: int64
```

Data Dictionary:


- age: age of the primary beneficiary - ratio (continuous variable)
- sex: sex of the beneficiary (male or female) - nominal (categorical variable)
- bmi: body mass index; a value derived from the mass and height of the beneficiary - interval (continuous variable)
- children: number of children covered by the insurance - ratio (discrete variable)
- smoker: whether the beneficiary smokes or not (yes or no) - nominal (categorical variable)
- region: residential area of the beneficiary in the U.S. - nominal (categorical variable)
- charges: individual medical costs billed by the insurance - ratio (continuous variable)

Step 4: Data Analytics


Create a plot of your choosing showing the data distributions for the following:

1. Age
2. Number of children covered by insurance
3. BMI
4. Smoking vs non-smoking
5. Charges

```
In [8]: # Write Step 4 Code Here:
fig, axes = plt.subplots(nrows=1, ncols=3,figsize=(18,6))
plotage = sns.histplot(data.age, color = 'lightblue', ax=axes[0])
plotchild = sns.histplot(data.children, color = 'yellow', ax=axes[1])
plotsmoker = sns.histplot(data.smoker, color = 'orange', ax=axes[2])
plt.show()
```




```
In [9]: fig, axes = plt.subplots(nrows=1, ncols=2,figsize=(18,6))
plot1 = sns.histplot(data.bmi, kde = False, color = 'lightgreen', ax=axes[0])
plot2 = sns.boxplot(data = data, x = 'bmi', color = 'lightgreen', ax=axes[1])
plt.show()
```



```
In [10]: fig, axes = plt.subplots(nrows=1, ncols=2,figsize=(18,6))
plot1 = sns.histplot(data.charges, kde = False, color = 'pink', ax=axes[0])
```

```
plot2 = sns.boxplot(data = data, x = 'charges', color = 'pink', ax=axes[1])
plt.show()
```



```
In [11]: data.describe()
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

Step 5: Use any method of your choosing to answer the following questions.

Are there any associates between smoking vs non-smoking and age?

```
In [12]: data.smoker = data.smoker.map(dict(yes=1, no=0))
data.sex = data.sex.map(dict(male=1,female=0))
data.head()
```

Out [12]:

	age	sex	bmi	children	smoker	region	charges
0	19	0	27.900	0	1	southwest	16884.92400
1	18	1	33.770	1	0	southeast	1725.55230
2	28	1	33.000	3	0	southeast	4449.46200
3	33	1	22.705	0	0	northwest	21984.47061
4	32	1	28.880	0	0	northwest	3866.85520

```
In [13]: data.corr()
```

Out [13]:

	age	sex	bmi	children	smoker	charges
age	1.000000	-0.020856	0.109272	0.042469	-0.025019	0.299008
sex	-0.020856	1.000000	0.046371	0.017163	0.076185	0.057292
bmi	0.109272	0.046371	1.000000	0.012759	0.003750	0.198341
children	0.042469	0.017163	0.012759	1.000000	0.007673	0.067998
smoker	-0.025019	0.076185	0.003750	0.007673	1.000000	0.787251
charges	0.299008	0.057292	0.198341	0.067998	0.787251	1.000000


The correlation between smoking vs non-smoking and age is -0.067428, no association.

Step 6: Perform an analysis of your choosing on the data to determine if any of the features within the data are predictive of cost.

In the interview we will ask you to explain how you approached this question.

```
In [14]: # Write Step 6 Code Here:
X = data[['age','sex','bmi','children','smoker','region']]
X = pd.get_dummies(data=X, drop_first=True)
X.head()
import statsmodels.api as sm
from scipy import stats
```

```
In [15]: # Log transform a single column
data['charges'] = np.log(data['charges'])
fig, axes = plt.subplots(nrows=1, ncols=2,figsize=(18,6))
plot1 = sns.histplot(data.charges, kde = False, color = 'pink', ax=axes[0])
plot2 = sns.boxplot(data = data, x = 'charges', color = 'pink', ax=axes[1])
plt.show()
```



```
In [16]: from sklearn.linear_model import LinearRegression
lmmodel = LinearRegression()
y = data['charges']
lmmodel.fit(X,y)
print('Intercept: \n', lmmodel.intercept_)
print('Coefficients: \n', lmmodel.coef_)

Intercept:
7.030558089523366
Coefficients:
[ 0.03458164 -0.07541644  0.01337482  0.10185685  1.55432279 -0.06378756
-0.15719675 -0.12895222]
```

```
In [17]: X2 = sm.add_constant(X)
est = sm.OLS(y, X2)
est2 = est.fit()
print(est2.summary())
```

OLS Regression Results						
=====						
Dep. Variable:	charges	R-squared:	0.768			
Model:	OLS	Adj. R-squared:	0.767			
Method:	Least Squares	F-statistic:	549.8			
Date:	Sun, 19 Feb 2023	Prob (F-statistic):	0.00			
Time:	16:27:59	Log-Likelihood:	-808.52			
No. Observations:	1338	AIC:	1635.			
Df Residuals:	1329	BIC:	1682.			
Df Model:	8					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[ 0.025	0.975]
-----						
const	7.0306	0.072	97.112	0.000	6.889	7.173
age	0.0346	0.001	39.655	0.000	0.033	0.036
sex	-0.0754	0.024	-3.091	0.002	-0.123	-0.028
bmi	0.0134	0.002	6.381	0.000	0.009	0.017
children	0.1019	0.010	10.085	0.000	0.082	0.122
smoker	1.5543	0.030	51.333	0.000	1.495	1.614
region_northwest	-0.0638	0.035	-1.827	0.068	-0.132	0.005
region_southeast	-0.1572	0.035	-4.481	0.000	-0.226	-0.088
region_southwest	-0.1290	0.035	-3.681	0.000	-0.198	-0.060
=====						
Omnibus:	463.882	Durbin-Watson:	2.046			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1673.760			
Skew:	1.679	Prob(JB):	0.00			
Kurtosis:	7.330	Cond. No.	311.			
=====						

Notes:  
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [18]: np.exp(lmmodel.intercept_)
```

Out[18]: 1130.6614444456202

```
In [19]: (np.exp(lmmodel.coef_)-1)*100
```

Out[19]: array([ 3.51865381, -7.26427867, 1.34646658, 10.72249595,  
 373.18809433, -6.17957129, -14.54640843, -12.09840377])


Since p-value of **region** is greater than 0.05, we would not reject the null hypothesis, so these region is not statistically significant.

Then we can conclude that **age, sex, bmi, the number of children covered by insurance, and whether the beneficiary smokes or not** are the features would within predictive analysis of medical cost.

- Holding all other features fixed, for a 1 year increase in Age, we expect to see about a 3.518% increase in the medical cost
- Holding all other features fixed, the medical cost will be 7.264% lower for MALE compare to FEMALE
- Holding all other features fixed, for a 1 unit increase in BMI, we expect to see about a 1.3468% increase in the medical cost
- Holding all other features fixed, for a 1 more Child, we expect to see about a 10.722% increase in the medical cost
- Holding all other features fixed, the medical cost will be 373.188% higher for SMOKER compare to NON-SMOKER

```
In [20]: from sklearn.model_selection import train_test_split
X2 = data[['age', 'sex', 'bmi', 'children', 'smoker']]
y = data['charges']
X2_train, X2_test, y_train, y_test = train_test_split(X2, y, test_size=0.2, random_state=101)
lmpredict = LinearRegression()
lmpredict.fit(X2_train,y_train)
predictions = lmpredict.predict(X2_test)
plt.scatter(y_test,predictions)
plt.xlabel('Y Test')
plt.ylabel('Predicted Y')
```

Out[20]: Text(0, 0.5, 'Predicted Y')



```
In [21]: from sklearn import metrics
print('MAE:', metrics.mean_absolute_error(y_test, predictions))
print('MSE:', metrics.mean_squared_error(y_test, predictions))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, predictions)))
```

MAE: 0.27135723968430475  
MSE: 0.17852875494422138  
RMSE: 0.4225266322307049

```
In [22]: sns.distplot((y_test-predictions),bins=50)
```

/Users/Cloris/opt/anaconda3/lib/python3.9/site-packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)

Out[22]: <AxesSubplot:xlabel='charges', ylabel='Density'>

